

```
!pip install git+https://github.com/RaRe-Technologies/gensim

Collecting git+https://github.com/RaRe-Technologies/gensim
  Cloning https://github.com/RaRe-Technologies/gensim to /tmp/pip-req-build-cmanzefg
    Running command git clone --filter=blob:none --quiet
https://github.com/RaRe-Technologies/gensim /tmp/pip-req-build-cmanzefg
      Resolved https://github.com/RaRe-Technologies/gensim to commit
8f815450ebc066825e2fbcd69669fe1fb0357143
  Installing build dependencies ... gets to build wheel ... etadata
(pyproject.toml) ... ent already satisfied: numpy>=1.18.5 in
/usr/local/lib/python3.12/dist-packages (from gensim==4.4.0a0.dev0)
(2.0.2)
Requirement already satisfied: scipy>=1.7.0 in
/usr/local/lib/python3.12/dist-packages (from gensim==4.4.0a0.dev0)
(1.16.2)
Requirement already satisfied: smart_open>=1.8.1 in
/usr/local/lib/python3.12/dist-packages (from gensim==4.4.0a0.dev0)
(7.3.1)
Requirement already satisfied: wrapt in
/usr/local/lib/python3.12/dist-packages (from smart_open>=1.8.1-
>gensim==4.4.0a0.dev0) (1.17.3)
Building wheels for collected packages: gensim
  Building wheel for gensim (pyproject.toml) ... : filename=gensim-
4.4.0a0.dev0-cp312-cp312-linux_x86_64.whl size=27080686
sha256=18820bf4c917b08f9a261e5fab2edf06126bb42c20d70c8d2301400c58779a1
4
  Stored in directory:
/tmp/pip-ephem-wheel-cache-bywpt9a8/wheels/77/a7/a3/cf0bad57a872edb85b
b0fe77d813d6b4b9f957a597a914587b
Successfully built gensim
Installing collected packages: gensim
Successfully installed gensim-4.4.0a0.dev0

import gensim
print("Gensim:", gensim.__version__)

Gensim: 4.3.3

!pip install umap-learn matplotlib scikit-learn plotly

Requirement already satisfied: umap-learn in
/usr/local/lib/python3.12/dist-packages (0.5.9.post2)
Requirement already satisfied: matplotlib in
/usr/local/lib/python3.12/dist-packages (3.10.0)
Requirement already satisfied: scikit-learn in
/usr/local/lib/python3.12/dist-packages (1.6.1)
Requirement already satisfied: plotly in
/usr/local/lib/python3.12/dist-packages (5.24.1)
```

```
Requirement already satisfied: numpy>=1.23 in
/usr/local/lib/python3.12/dist-packages (from umap-learn) (2.0.2)
Requirement already satisfied: scipy>=1.3.1 in
/usr/local/lib/python3.12/dist-packages (from umap-learn) (1.16.2)
Requirement already satisfied: numba>=0.51.2 in
/usr/local/lib/python3.12/dist-packages (from umap-learn) (0.60.0)
Requirement already satisfied: pynndescent>=0.5 in
/usr/local/lib/python3.12/dist-packages (from umap-learn) (0.5.13)
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-
packages (from umap-learn) (4.67.1)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.12/dist-packages (from matplotlib) (1.3.3)
Requirement already satisfied: cycler>=0.10 in
/usr/local/lib/python3.12/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.12/dist-packages (from matplotlib) (4.60.1)
Requirement already satisfied: kiwisolver>=1.3.1 in
/usr/local/lib/python3.12/dist-packages (from matplotlib) (1.4.9)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.12/dist-packages (from matplotlib) (25.0)
Requirement already satisfied: pillow>=8 in
/usr/local/lib/python3.12/dist-packages (from matplotlib) (11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.12/dist-packages (from matplotlib) (3.2.5)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.12/dist-packages (from matplotlib)
(2.9.0.post0)
Requirement already satisfied: joblib>=1.2.0 in
/usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.5.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/usr/local/lib/python3.12/dist-packages (from scikit-learn) (3.6.0)
Requirement already satisfied: tenacity>=6.2.0 in
/usr/local/lib/python3.12/dist-packages (from plotly) (8.5.0)
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in
/usr/local/lib/python3.12/dist-packages (from numba>=0.51.2->umap-
learn) (0.43.0)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.7-
>matplotlib) (1.17.0)

import random
import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE
import umap
import gensim
import gensim.downloader as api
from gensim.models import Word2Vec, FastText
from sklearn.metrics.pairwise import cosine_similarity
```

```
import plotly.express as px

%matplotlib inline
plt.rcParams['figure.figsize'] = (10,7)

print("Gensim:", gensim.__version__)
Gensim: 4.3.3

from itertools import islice

# Tải corpus text8 (generator)
corpus = api.load("text8")

# Lấy 50k câu đầu tiên để huấn luyện nhanh
sentences = list(islice(corpus, 50000))
print("Số câu dùng để huấn luyện:", len(sentences))

Số câu dùng để huấn luyện: 1701

# Tải GloVe pretrained
glove = api.load("glove-wiki-gigaword-100")
print("GloVe vocab size:", len(glove.key_to_index))

[=====] 100.0%
128.1/128.1MB downloaded
GloVe vocab size: 400000

# Gia tăng dữ liệu xuông 10k câu
from itertools import islice
sentences_small = list(islice(api.load("text8"), 10000))

# Huấn luyện Word2Vec nhanh
w2v_model = Word2Vec(sentences=sentences_small, vector_size=50,
window=5, min_count=5, sg=1, epochs=3)
print("Word2Vec vocab size:", len(w2v_model.wv.index_to_key))

# Huấn luyện FastText nhanh
ft_model = FastText(sentences=sentences_small, vector_size=50,
window=5, min_count=5, epochs=3)
print("FastText vocab size:", len(ft_model.wv.index_to_key))

Word2Vec vocab size: 71290

-----
-----
KeyboardInterrupt                                     Traceback (most recent call
last)
/tmp/ipython-input-65606408.py in <cell line: 0>()
      8
      9 # Huấn luyện FastText nhanh
--> 10 ft_model = FastText(sentences=sentences_small, vector_size=50,
```

```

window=5, min_count=5, epochs=3)
    11 print("FastText vocab size:", len(ft_model.wv.index_to_key))

/usr/local/lib/python3.12/dist-packages/gensim/models/fasttext.py in
__init__(self, sentences, corpus_file, sg, hs, vector_size, alpha,
window, min_count, max_vocab_size, word_ngrams, sample, seed, workers,
min_alpha, negative, ns_exponent, cbow_mean, hashfxn, epochs,
null_word, min_n, max_n, sorted_vocab, bucket, trim_rule, batch_words,
callbacks, max_final_vocab, shrink_windows)
    433         self.wv.vectors_ngrams_lockf = ones(1, dtype=REAL)
    434
--> 435         super(FastText, self).__init__(
    436             sentences=sentences, corpus_file=corpus_file,
workers=workers, vector_size=vector_size, epochs=epochs,
    437             callbacks=callbacks, batch_words=batch_words,
trim_rule=trim_rule, sg=sg, alpha=alpha, window=window,

/usr/local/lib/python3.12/dist-packages/gensim/models/word2vec.py in
__init__(self, sentences, corpus_file, vector_size, alpha, window,
min_count, max_vocab_size, sample, seed, workers, min_alpha, sg, hs,
negative, ns_exponent, cbow_mean, hashfxn, epochs, null_word,
trim_rule, sorted_vocab, batch_words, compute_loss, callbacks,
comment, max_final_vocab, shrink_windows)
    428
self._check_corpus_sanity(corpus_iterable=corpus_iterable,
corpus_file=corpus_file, passes=(epochs + 1))
    429         self.build_vocab(corpus_iterable=corpus_iterable,
corpus_file=corpus_file, trim_rule=trim_rule)
--> 430         self.train(
    431             corpus_iterable=corpus_iterable,
corpus_file=corpus_file, total_examples=self.corpus_count,
    432             total_words=self.corpus_total_words,
epochs=self.epochs, start_alpha=self.alpha,
    433             report_delay=report_delay)

/usr/local/lib/python3.12/dist-packages/gensim/models/word2vec.py in
train(self, corpus_iterable, corpus_file, total_examples, total_words,
epochs, start_alpha, end_alpha, word_count, queue_factor,
report_delay, compute_loss, callbacks, **kwargs)
    1071
    1072         if corpus_iterable is not None:
-> 1073             trained_word_count_epoch,
raw_word_count_epoch, job_tally_epoch = self._train_epoch(
    1074                 corpus_iterable, cur_epoch=cur_epoch,
total_examples=total_examples,
    1075                 total_words=total_words,
queue_factor=queue_factor, report_delay=report_delay,
    1076

/usr/local/lib/python3.12/dist-packages/gensim/models/word2vec.py in
_train_epoch(self, data_iterable, cur_epoch, total_examples,
total_words, queue_factor, report_delay, callbacks)

```

```

1432             thread.start()
1433
-> 1434         trained_word_count, raw_word_count, job_tally =
self._log_epoch_progress(
    1435             progress_queue, job_queue, cur_epoch=cur_epoch,
total_examples=total_examples,
    1436             total_words=total_words,
report_delay=report_delay, is_corpus_file_mode=False,

```

/usr/local/lib/python3.12/dist-packages/gensim/models/word2vec.py in \_log\_epoch\_progress(self, progress\_queue, job\_queue, cur\_epoch, total\_examples, total\_words, report\_delay, is\_corpus\_file\_mode)

```

1287
1288         while unfinished_worker_count > 0:
-> 1289             report = progress_queue.get() # blocks if workers
too slow
1290             if report is None: # a thread reporting that it
finished
1291                 unfinished_worker_count -= 1

```

/usr/lib/python3.12/queue.py in get(self, block, timeout)

```

169             elif timeout is None:
170                 while not self._qsize():
--> 171                     self.not_empty.wait()
172             elif timeout < 0:
173                 raise ValueError("'timeout' must be a non-
negative number")

```

/usr/lib/python3.12/threading.py in wait(self, timeout)

```

353         try: # restore state no matter what (e.g.,
KeyboardInterrupt)
    354             if timeout is None:
--> 355                 waiter.acquire()
    356             gotit = True
    357         else:

```

KeyboardInterrupt:

```

from itertools import islice

# Lấy ít dữ liệu hơn cho nhanh
small_sentences = list(islice(api.load("text8"), 2000)) # chỉ 2000
câu

# FastText demo nhẹ
ft_model = FastText(sentences=small_sentences, vector_size=30,
window=3, min_count=2, epochs=2)
print("FastText vocab size:", len(ft_model.wv.index_to_key))

FastText vocab size: 135335

```

```

words_to_plot = [
    "king", "queen", "man", "woman", "prince", "princess",
    "apple", "orange", "banana", "fruit",
    "paris", "france", "london", "england", "germany",
    "cat", "dog", "horse", "cow", "pig",
    "run", "running", "walk", "swim", "eat"
]

def filter_vocab(word_list, model):
    return [w for w in word_list if w in model.wv.index_to_key]

w_w2v = filter_vocab(words_to_plot, w2v_model)
w_ft = filter_vocab(words_to_plot, ft_model)
w_glove = [w for w in words_to_plot if w in glove.key_to_index]

print("Word2Vec words:", w_w2v)
print("FastText words:", w_ft)
print("GloVe words:", w_glove)

Word2Vec words: ['king', 'queen', 'man', 'woman', 'prince',
'princess', 'apple', 'orange', 'banana', 'fruit', 'paris', 'france',
'london', 'england', 'germany', 'cat', 'dog', 'horse', 'cow', 'pig',
'run', 'running', 'walk', 'swim', 'eat']
FastText words: ['king', 'queen', 'man', 'woman', 'prince',
'princess', 'apple', 'orange', 'banana', 'fruit', 'paris', 'france',
'london', 'england', 'germany', 'cat', 'dog', 'horse', 'cow', 'pig',
'run', 'running', 'walk', 'swim', 'eat']
GloVe words: ['king', 'queen', 'man', 'woman', 'prince', 'princess',
'apple', 'orange', 'banana', 'fruit', 'paris', 'france', 'london',
'england', 'germany', 'cat', 'dog', 'horse', 'cow', 'pig', 'run',
'running', 'walk', 'swim', 'eat']

import plotly.express as px
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE
import umap

def get_vectors(model, words):
    return np.array([model.wv[w] for w in words])

def plot_2d(points, labels, title="2D projection"):
    plt.figure(figsize=(10,8))
    plt.scatter(points[:,0], points[:,1])
    for i, txt in enumerate(labels):
        plt.annotate(txt, (points[i,0], points[i,1]))
    plt.title(title)
    plt.grid(True)
    plt.show()

def plot_3d(points, labels, title="3D projection"):

```

```

    fig = px.scatter_3d(x=points[:,0], y=points[:,1], z=points[:,2],
text=labels, title=title)
    fig.update_traces(marker=dict(size=4))
    fig.show()

def reduce_and_plot_2d(model, words, method="pca"):
    vecs = get_vectors(model, words)
    if method == "pca":
        pts = PCA(n_components=2).fit_transform(vecs)
        plot_2d(pts, words, f"{type(model).__name__} - PCA 2D")
    elif method == "tsne":
        # chọn perplexity nhỏ hơn số từ
        perplexity = min(10, len(words)-1)
        pts = TSNE(n_components=2, random_state=42, init="pca",
perplexity=perplexity).fit_transform(vecs)
        plot_2d(pts, words, f"{type(model).__name__} - t-SNE 2D")
    elif method == "umap":
        pts = umap.UMAP(n_components=2,
random_state=42).fit_transform(vecs)
        plot_2d(pts, words, f"{type(model).__name__} - UMAP 2D")
def reduce_and_plot_3d(model, words, method="umap"):
    vecs = get_vectors(model, words)
    if method == "umap":
        pts = umap.UMAP(n_components=3,
random_state=42).fit_transform(vecs)
    else:
        pts = PCA(n_components=3).fit_transform(vecs)
    plot_3d(pts, words, f"{type(model).__name__} - {method.upper()} 3D")

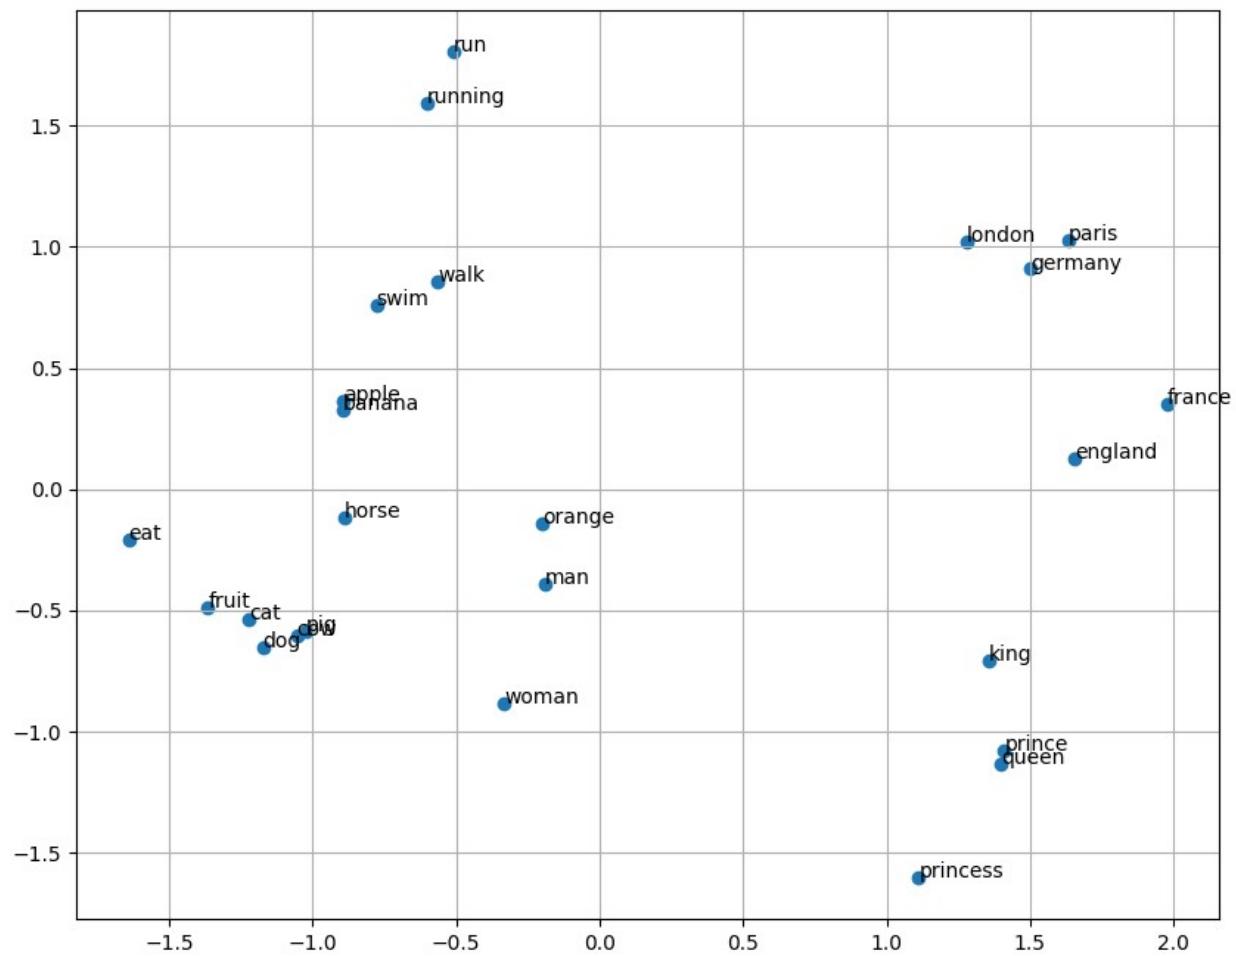
# Word2Vec
reduce_and_plot_2d(w2v_model, w_w2v, "pca")
reduce_and_plot_2d(w2v_model, w_w2v, "tsne")
reduce_and_plot_3d(w2v_model, w_w2v, "umap")

# FastText
reduce_and_plot_2d(ft_model, w_ft, "pca")
reduce_and_plot_2d(ft_model, w_ft, "umap")

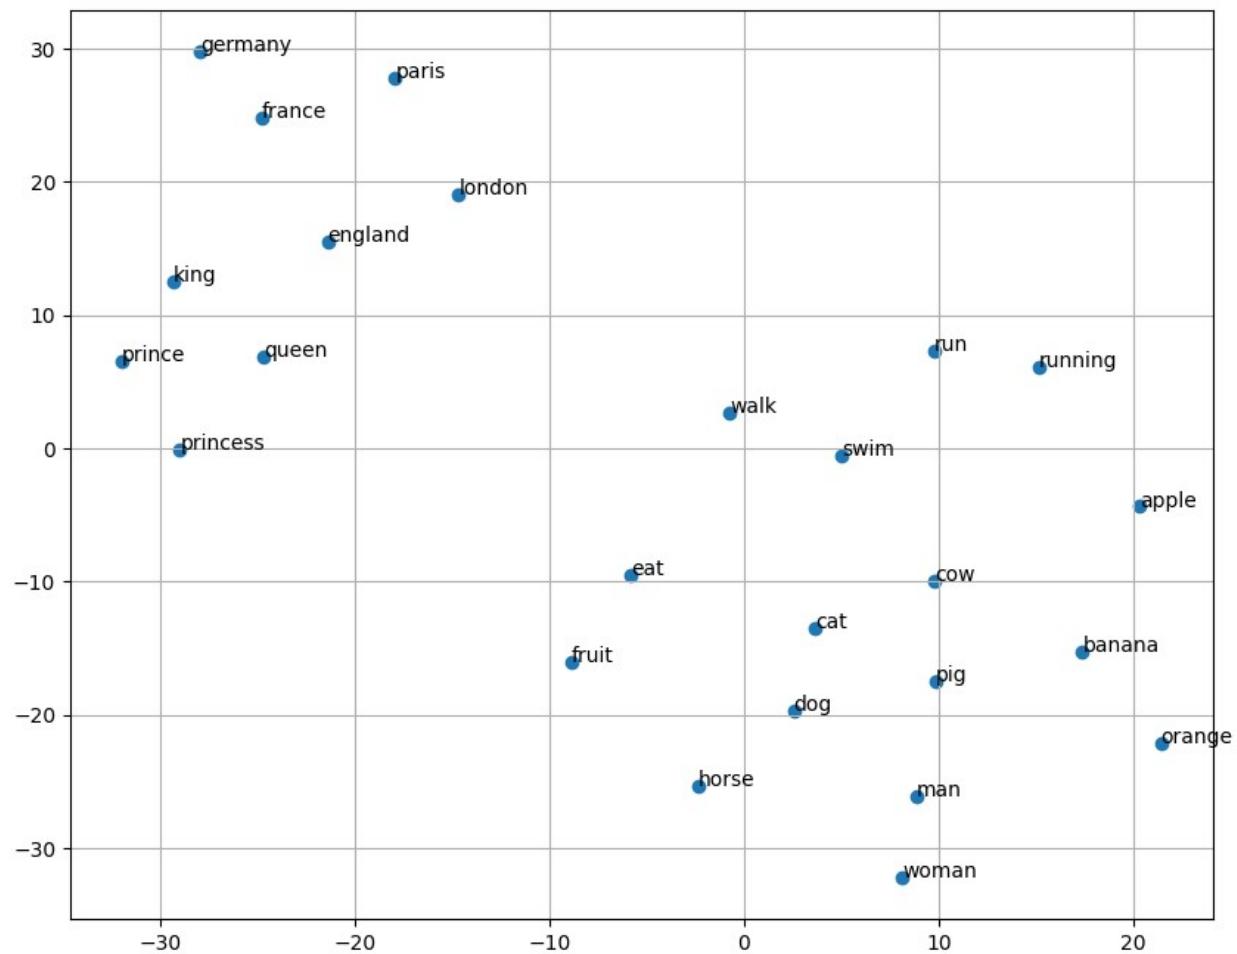
# GloVe
vecs_glove = np.array([glove[w] for w in w_glove])
pts_glove = PCA(n_components=2).fit_transform(vecs_glove)
plot_2d(pts_glove, w_glove, "GloVe - PCA 2D")

```

Word2Vec - PCA 2D

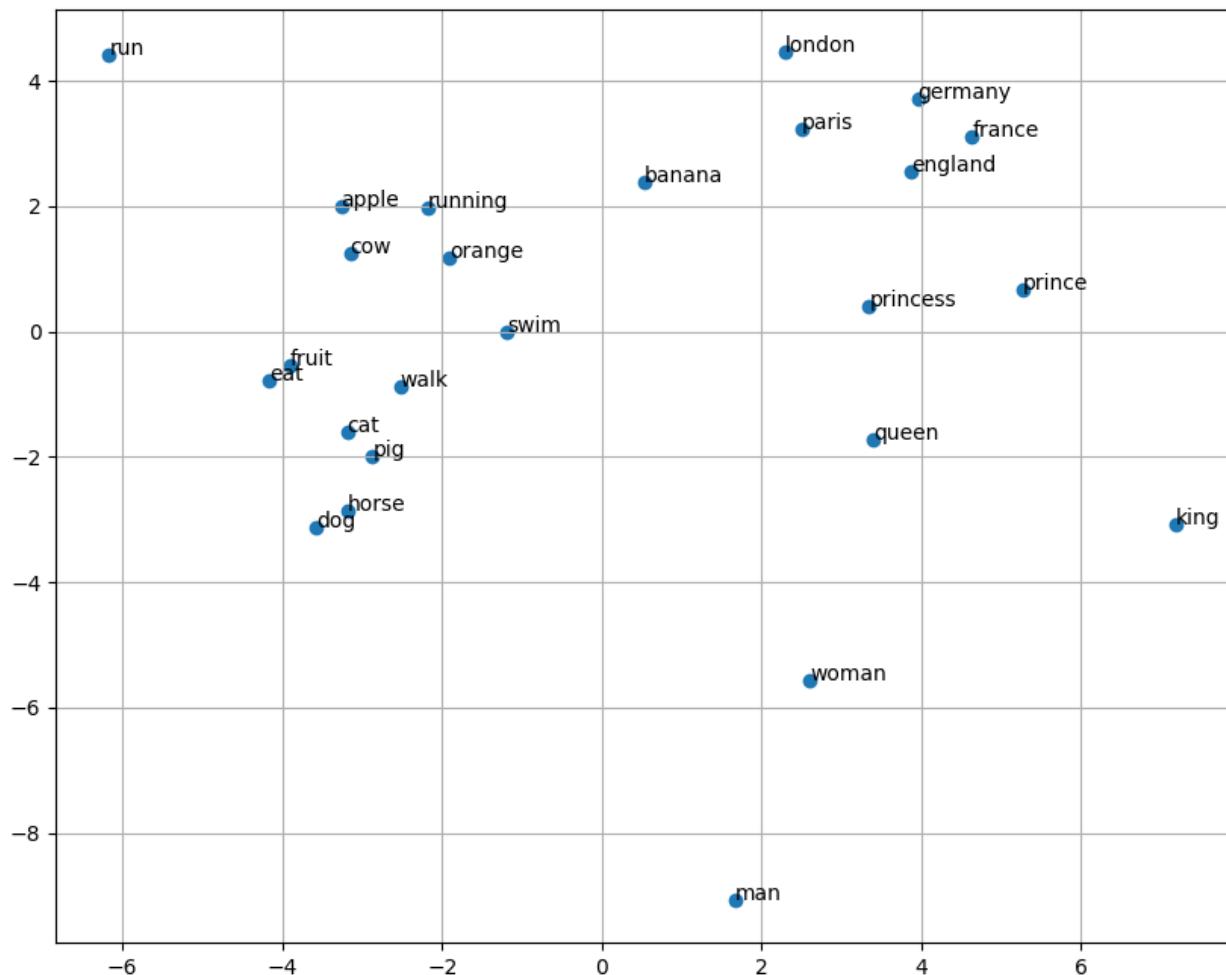


Word2Vec - t-SNE 2D



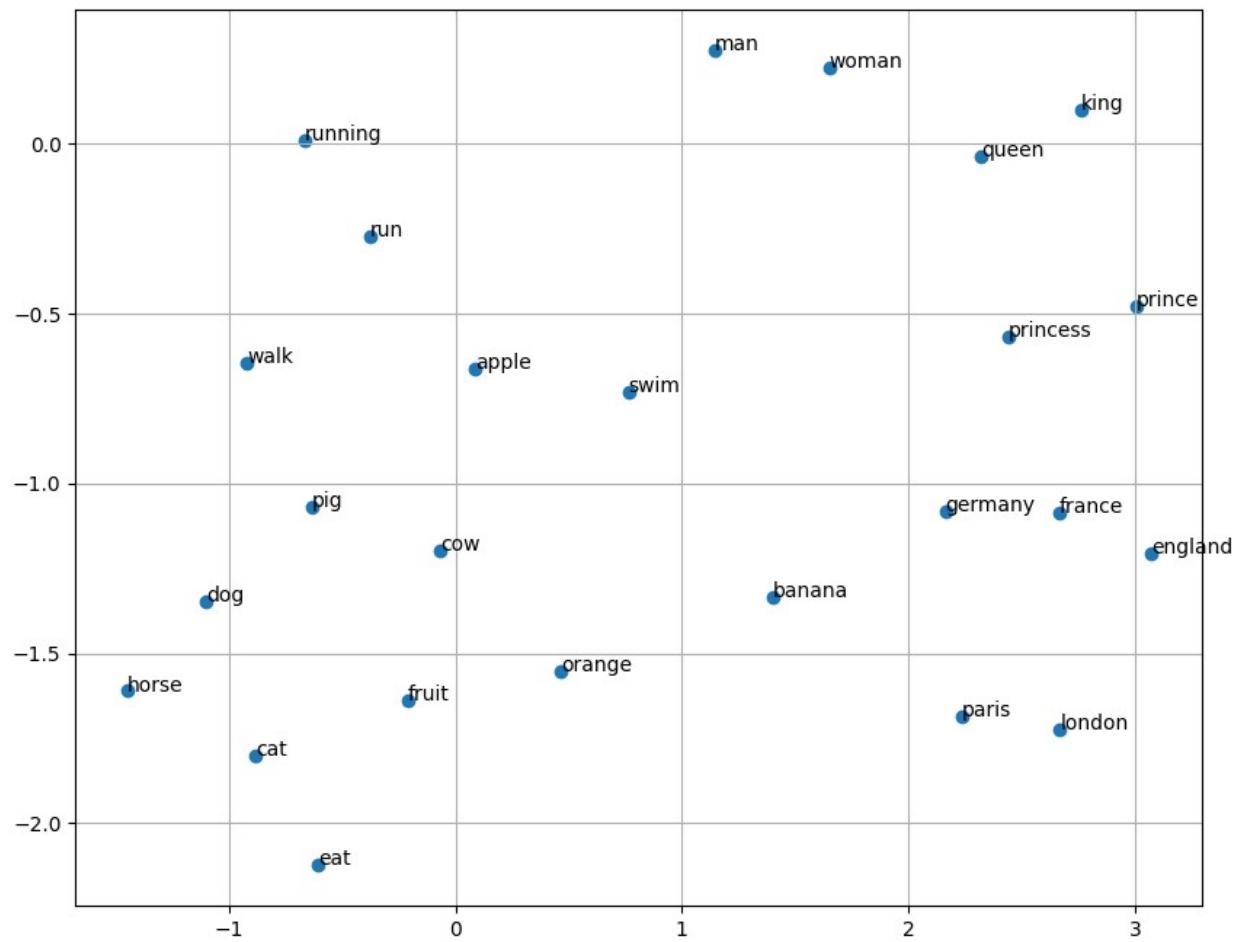
```
/usr/local/lib/python3.12/dist-packages/umap/umap_.py:1952:  
UserWarning: n_jobs value 1 overridden to 1 by setting random_state.  
Use no seed for parallelism.  
warn(
```

FastText - PCA 2D

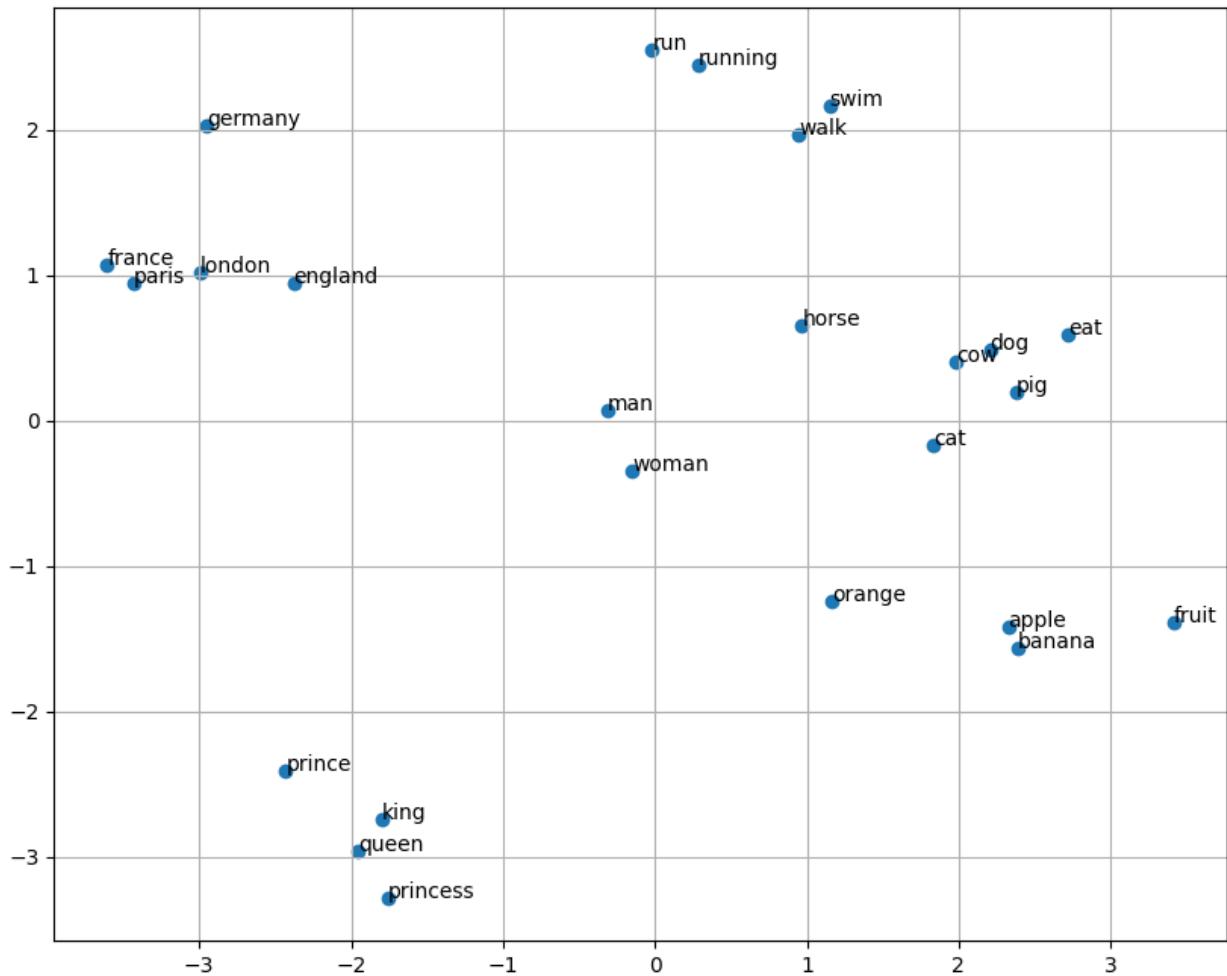


```
/usr/local/lib/python3.12/dist-packages/umap/umap_.py:1952:  
UserWarning:  
n_jobs value 1 overridden to 1 by setting random_state. Use no seed  
for parallelism.
```

FastText - UMAP 2D



GloVe - PCA 2D



```

from sklearn.metrics.pairwise import cosine_similarity

def top_k_similar(model, word, k=5):
    try:
        return model.wv.most_similar(word, topn=k)
    except:
        if word not in model.key_to_index:
            return f"'{word}' không có trong vocab"
        vec = model[word].reshape(1,-1)
        all_words = list(model.key_to_index.keys())
        all_vecs = np.array([model[w] for w in all_words])
        sims = cosine_similarity(vec, all_vecs)[0]
        idx = np.argsort(sims)[-k-1:-1][::-1]
        return [(all_words[i], float(sims[i])) for i in idx]

# Ví dụ
print("Word2Vec:", top_k_similar(w2v_model, "king"))
print("FastText:", top_k_similar(ft_model, "king"))
print("GloVe:", top_k_similar(glove, "king"))

```

```
Word2Vec: [('prince', 0.8788297772407532), ('olaf', 0.8643296360969543), ('valdemar', 0.864129900932312), ('queen', 0.8618401885032654), ('pretender', 0.8617712259292603)]  
FastText: [('kingswear', 0.9054992198944092), ('kingpin', 0.8961842656135559), ('kingship', 0.8770096302032471), ('kief', 0.8647317886352539), ('vingt', 0.8641076683998108)]  
GloVe: [('prince', 0.7682329416275024), ('queen', 0.7507690191268921), ('son', 0.7020888924598694), ('brother', 0.6985775232315063), ('monarch', 0.6977890729904175)]
```

## Nhận xét cá nhân

- PCA chạy nhanh, nhưng cụm từ hơi lẩn; t-SNE và UMAP cho cụm rõ ràng hơn.
- Word2Vec học được quan hệ ngữ nghĩa (king–queen, man–woman) khá rõ.
- FastText có lợi thế với từ ngoài từ điển hoặc biến thể (thêm tiền tố/hậu tố).
- GloVe pretrained cho kết quả tự nhiên và ý nghĩa hơn vì được huấn luyện trên tập dữ liệu lớn.
- Từ đây có thể thấy: trong thực tế, để đạt chất lượng cao nên dùng mô hình pretrained; còn huấn luyện từ đầu chỉ phù hợp demo hoặc dữ liệu chuyên biệt.